

Analysis of Floating-Point Matrix Multiplication Computed via Integer Arithmetic

Mantas Mikaitis

Joint work with A. Abdelfattah, J. Dongarra, M. Fasi, and F. Tisseur

School of Computer Science, University of Leeds, Leeds, UK


27th Conference of the International Linear Algebra Society (ILAS 2026)

Session: Approximate Computing in Numerical Linear Algebra

Virginia Tech, Blacksburg, VA, US

19 May, 2026

Fast low-precision matrix mult on the TOP500



—▲— 16-bit FP —●— 19-bit FP —○— 8-bit FP —×— int matrix mult.

Devices counted: V100, A100, H100, MI210, MI250X, MI300X, Intel Data Center GPU, from <https://www.top500.org>.

NVIDIA HGX B200 throughputs (OPS/s)

int8 (0.45×10^{16}) fp16 (0.22×10^{16}) fp64 (0.4×10^{14}).


What is the “Ozaki scheme”?

In a nutshell

Split matrices into several “slices”, compute many exact matrix multiplies, sum the products $\rightarrow$ more accurate than one $m \times k \times n$ product.

- [Ozaki, Ogita, Oishi, Rump, 2012] accurate fp64 mat mul by splitting the inputs and performing several exact fp64 mat mul + sum:
 $AB = \sum \text{fl}(A^{(r)} B^{(s)}).$
- [Ozaki, Ogita, Oishi, Rump, 2013] improve the algorithm by exploiting *last non-zero bit* to reduce the number of splits/products.
- [Mukunoki, Ozaki, Ogita, Imamura, 2020] used 16-bit input, 32-bit output FP tensor cores for the above.
- [Ootomo, Ozaki, Yokota, 2024] used 8-bit input, 32-bit output **integer** tensor cores.
- [Mukunoki, 2026] uses an 8-bit FP tensor core.
- [Schwarz et al., 2026] tackles optimal slicing to 8 bit INT.

Mukunoki et al. 2020 illustration of the Ozaki scheme



Using integer tensor cores

A hardware FP64 FMA is multiplying 53-bit integers and adding 106-bit integers—FP64 fundamentally is computed through integers.

Similarly, **we can use integer tensor cores for FP64 mat-mul computations.**

Take $A \in \mathbb{F}^{m \times k}$ and $B \in \mathbb{F}^{k \times n}$. Goal: approximate AB by integer tensor cores. Typically $\mathbb{F}$ is double prec.

Assume A and B have no infinities, NaNs, or -0 and the computation cannot produce inf and NaN.

Step 1: Denormalize floating-point data to align exponents

- 1 For each row i of A , compute a scale factor

$$\alpha_i = 2^{\lfloor \log_2 M_i \rfloor} \times 2, \quad M_i = \max_{1 \leq j \leq k} |a_{ij}|, \quad 1 \leq i \leq m,$$

- 2 Shift each a_{ij} significand right, until exponents $\log_2(\alpha_{ij})$.
- 3 This guarantees each element scaled down by α_i is in $[0, 1)$.
- 4 And $0.5 \leq M_i/\alpha_i < 1$.
- 5 Similar for B , **column-wise**.

Block floating-point arithmetic


Note that this is block floating-point representation of [\[Wilkinson, 1963\]](#).

Ozaki scheme: each row of A and column of B is a block with its own scale factor.

Utilising 8-bit integer tensor cores [Ootomo et al. 2024]

Step 2: Slicing

shared-place splitting (Ozaki scheme)



Scaling and slicing: example on small vectors

$$A = \begin{bmatrix} 1.5625 & 8 & -3.6875 \end{bmatrix}, \quad B = \begin{bmatrix} 1.3828125 \\ -7.625 \\ 3.625 \end{bmatrix}$$

Scale factors $\alpha = 2^4$ and $\beta = 2^3$.

Example set up: FP precision 8 bits, 4 slices, integer: 3 bits and a sign

$$\begin{aligned} \begin{bmatrix} 2^0 \cdot 1.1001000 \\ 2^3 \cdot 1.0000000 \\ -2^1 \cdot 1.1101100 \end{bmatrix} &\Rightarrow 2^4 \cdot \begin{bmatrix} \emptyset.000 \ 110 \ 010 \ 000 \\ \emptyset.100 \ 000 \ 000 \ 000 \\ -\emptyset.001 \ 110 \ 110 \ 000 \end{bmatrix} \Rightarrow 2^1 \cdot \begin{bmatrix} 000 \\ 100 \\ -001 \end{bmatrix} + 2^{-2} \cdot \begin{bmatrix} 110 \\ 000 \\ -110 \end{bmatrix} + 2^{-5} \cdot \begin{bmatrix} 010 \\ 000 \\ -110 \end{bmatrix} + 2^{-8} \cdot \begin{bmatrix} 000 \\ 000 \\ 000 \end{bmatrix} \\ A^T &\quad \text{Block fixed-point} \quad A_{(1)}^T \quad A_{(2)}^T \quad A_{(3)}^T \quad A_{(4)}^T \\ \\ \begin{bmatrix} 2^0 \cdot 1.0110001 \\ -2^2 \cdot 1.1110100 \\ 2^1 \cdot 1.1101000 \end{bmatrix} &\Rightarrow 2^3 \cdot \begin{bmatrix} \emptyset.001 \ 011 \ 000 \ 100 \\ -\emptyset.111 \ 101 \ 000 \ 000 \\ \emptyset.011 \ 101 \ 000 \ 000 \end{bmatrix} \Rightarrow 2^0 \cdot \begin{bmatrix} 001 \\ -111 \\ 011 \end{bmatrix} + 2^{-3} \cdot \begin{bmatrix} 011 \\ -101 \\ 101 \end{bmatrix} + 2^{-6} \cdot \begin{bmatrix} 000 \\ 000 \\ 000 \end{bmatrix} + 2^{-9} \cdot \begin{bmatrix} 100 \\ 000 \\ 000 \end{bmatrix} \\ B &\quad \text{Block fixed-point} \quad B^{(1)} \quad B^{(2)} \quad B^{(3)} \quad B^{(4)} \end{aligned}$$

We now have:

$$A \approx \text{diag}(\alpha) \sum_{\ell=1}^{s_A} 2^{-\ell t} A_{(\ell)}, \quad B \approx \sum_{h=1}^{s_B} 2^{-ht} B^{(h)} \text{diag}(\beta).$$

Here:

- s_A and s_B are the number of slices of BFP A and B .
- $A_{(\ell)}$ is the ℓ -th slice of A .
- $B^{(h)}$ is the h -th slice of B .
- t is the precision of integers in $A_{(\ell)}$ and $B^{(h)}$.

Step 3: Integer products

$$\begin{aligned}
 C = AB &\approx \left(\text{diag}(\alpha) \sum_{\ell=1}^{s_A} 2^{-\ell t} A_{(\ell)} \right) \left(\sum_{h=1}^{s_B} 2^{-ht} B^{(h)} \text{diag}(\beta) \right) \\
 &= \alpha \beta^T \circ \sum_{\ell=1}^{s_A} \sum_{h=1}^{s_B} 2^{-(\ell+h)t} \underline{A_{(\ell)} B^{(h)}}.
 \end{aligned}$$

Step 4: Summation and rescaling (back in FP64)

$$\underline{C \approx \alpha \beta^T \circ \sum_{\ell=1}^{s_A} \sum_{h=1}^{s_B} 2^{-(\ell+h)t} A_{(\ell)} B^{(h)}}.$$

Rounding error analysis 1: slicing

Error induced in slicing

Slicing causes the error in the virtual block floating-point precision, by taking a set number of slices s_A and s_B . The precision to the right of the binary point in BFP is $s_A t$ bits—**the more slices we take, the higher the precision of the input is**. However, BFP causes errors relative to the max value in a block.

For slicing A (analogous for B) we have the **absolute error**:

$$A = \Delta A + \text{diag}(\alpha) \sum_{\ell=1}^{s_A} 2^{-\ell t} A_{(\ell)}, \quad |\delta a_{ij}| < \alpha_i 2^{-s_A t},$$

and the relative error (component wise)

$$\frac{|\delta a_{ij}|}{|a_{ij}|} < \frac{\alpha_i}{|a_{ij}|} 2^{-s_A t} \leq \frac{2 \max_j |a_{ij}|}{\min_j |a_{ij}|} 2^{-s_A t} \leq \kappa_A 2^{-s_A t}, \quad \kappa_A := 2 \max_i \frac{\max_j |a_{ij}|}{\min_j |a_{ij}|}$$

Rounding error analysis 2: product summation

Let $\hat{C}$ be an FP approximation for $\alpha\beta^T \circ \sum_{\ell=1}^{s_A} \sum_{h=1}^{s_B} 2^{-(\ell+h)t} A_{(\ell)} B^{(h)}$.

We arrived at the bound for $\hat{C}$:

$$|\hat{C} - C| \lesssim \left(2^{-s_A t} \kappa_A + 2^{-s_B t} \kappa_B + (s_A s_B - 1) 2^{-p} \right) |A| |B|.$$

where p is the precision of the FP format used for accumulation (e.g. FP64).

Key take aways

- Large error if κ get large (large ranges of A and/or B).
- Taking more slices can help: inc. s_A or s_B . But more expensive.
- If $s_A t$ and $s_B t$ are large enough, 2^{-p} (accumulation error) dominates.

Experiment 1: 2-element vectors

As a minimal example, we consider the computation of the inner product $a^T b$, where

$$a = \begin{bmatrix} 2^{-\varphi} x \\ 1 \end{bmatrix}, \quad b = \begin{bmatrix} 2^{\varphi} y \\ 1 \end{bmatrix}, \quad x, y \sim \mathcal{N}(0, 1).$$


We plot:

$$\frac{|\hat{c} - c|}{|c|},$$

where

- $\hat{c}$ is computed with Ozaki scheme on 8-bit to 32-bit int tensor core.
- c is a reference solution computed using the MATLAB Symbolic Toolbox with 32 decimal digits of accuracy.

Experiment 1: 2-element vectors



The x-axis denotes the number of slices and the y-axis controls the wideness of the gap between the min and max exponents.

Key takeaways

- For $\varphi = 0$ about 7 slices are sufficient for FP64 accur.
- For $\varphi = 100$ about 20 slices are needed.

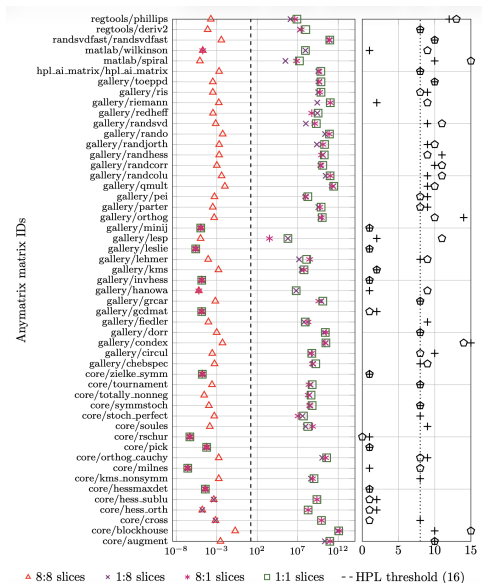
Experiment 2: Solving $Ax = b$ with block LU

- Take $A \in \mathbb{R}^{500 \times 500}$, one of the matrices from the Anymatrix collection [Higham and Mikaitis, 2021].
- $b \in \mathbb{R}^{500}$ with entries $[0, 1)$ from unif. distr.
- Use block LU with partial pivoting and block size 10.
- All mat-mul computed with the Ozaki scheme inside the LU iteration.

HPL error measure (used for <https://top500.org>)


Report: $\frac{\|A\hat{x} - b\|_\infty}{2^u(\|A\|_\infty \|\hat{x}\|_\infty + \|b\|_\infty)n}$ (HPL pass threshold < 16). Here $u = 2^{-53}$ for FP64.

Experiment 2: Solving $Ax = b$ with block LU



Experiment 2: Solving $Ax = b$ with block LU


Error $\frac{\|A\hat{x} - b\|_\infty}{2u(\|A\|_\infty \|\hat{x}\|_\infty + \|b\|_\infty)n}$ in solving $Ax = b$ for A an Anymatrix matrix core/blockhouse with $n = 500$.



HPL threshold 16 passed for configs in dotted square.

Experiment 2: Solving $Ax = b$ with block LU

Error $\frac{\|A\hat{x} - b\|_\infty}{2u(\|A\|_\infty\|\hat{x}\|_\infty + \|b\|_\infty)n}$ in solving $Ax = b$ for A an Anymatrix matrix core/cross with $n = 500$.



HPL threshold 16 passed for configs in dotted square.

Summary

- Low-precision matrix multipliers are being used for general-purpose computation.
- We have
 - Built error analysis to check some intuitive predictions about the Ozaki scheme.
 - Experimentation to show where the scheme works or not.
 - Unbalanced slicing: future work.

Preprint

A. bdelfattah, J. Dongarra, M. Fasi, M. Mikaitis, and F. Tisseur. *Analysis of Floating-Point Matrix Multiplication Computed via Integer Arithmetic*. **Preprint, arXiv:2506.11277 [math.NA], Accepted for SIAM SISC.**
Jun. 2025.

Slides at <http://mmikaitis.github.io/conferences>

References I



J. Wilkinson

Rounding Errors in Algebraic Processes

Her Majesty's Stationery Office, 1963 (Reprinted by SIAM in 2023)



K. Ozaki, T. Ogita, S. Oishi, S. M. Rump

Error-free transformations of matrix multiplication by using fast routines of matrix multiplication and its applications

Numer. Alg., 59. 2012.



D. Mukunoki, K. Ozaki, T. Ogita, T. Imamura

DGEMM Using Tensor Cores, and Its Accurate and Reproducible Versions

LNCS 12151. 2020.



N. J. Higham and M. Mikaitis

Anymatrix: An Extensible MATLAB Matrix Collection

Numer. Algorithms, 90. 2021.

References II



H. Ootomo, K. Ozaki and R. Yokota

DGEMM on integer matrix multiplication unit
Int. J. High Perf. Comput. Appl., 38. 2024.



D. Mukunoki

DGEMM using FP64 Arithmetic Emulation and FP8 Tensor Cores
with Ozaki Scheme
SCA/HPCAsia 2026 Workshops, Osaka, Japan



A. Schwarz, A. Anders, C. Brower, H. Bayraktar, J. Gunnels, and K. Clark

Guaranteed DGEMM Accuracy While Using Reduced Precision Tensor Cores Through Extensions of the Ozaki Scheme
SCA/HPCAsia 2026 Workshops, Osaka, Japan

(Extra slides) Experiment 3: matrices with increasing dyn. range

We take the matrices $A \in \mathbb{R}^{10 \times k}$ and $B \in \mathbb{R}^{k \times 10}$ with

$$A_{ij} = a_{ij} e^{\varphi x_{ij}}, \quad B_{ij} = b_{ij} e^{\varphi y_{ij}}, \quad a_{ij}, b_{ij} \sim \mathcal{U}(-0.5, 0.5), \quad x_{ij}, y_{ij} \sim \mathcal{N}(0, 1),$$

The parameter φ allows us to control the exponent range of the elements in A and B .

We plot the max forw. rel. error (following [\[Ootomo, Ozaki, Yokota, 2024\]](#)):

$$\max_{i,j} \frac{|\hat{c}_{ij} - c_{ij}|}{|c_{ij}|},$$

(Extra slides) Experiment 3: matrices with increasing dyn. range

